

# **ModelarDB: Modular Model-Based Time Series Management with Spark and Cassandra**

**Søren Kejser Jensen (Aalborg University)**

**Torben Bach Pedersen (Aalborg University)**

**Christian Thomsen (Aalborg University)**


Center for Data-intensive Systems

# Motivation



- Large industrial systems produce big amounts of **high quality sensor data**
- Data is collected as **regular time series** with **only a few gaps** without values
- High frequency could benefit analysis but:
  - High frequency data requires very fast ingestion
  - High query processing time limits use of historical data
  - Unfeasible high amounts of storage are required
- Storage of real-life wind turbine data:

Storage Method	Size (GiB)	Storage Method	Size (GiB)
PostgreSQL 10.1	782.87	CSV Files	582.68
RDBMS-X - Row	367.89	Apache Parquet Files	106.94
RDBMS-X - Column	166.83	Apache ORC Files	13.50
InfluxDB 1.4.2 - Tags	4.33	Apache Cassandra 3.9	111.89
InfluxDB 1.4.2 - Measurements	4.33	<i>ModelarDB</i>	2.41 - 2.84



- Currently **simple aggregates** are stored with **outliers and fluctuations lost!**

# Model-based storage of time series



- A **model** is a **lossy** or **lossless** representation of a time series
- E.g., a linear function reduces the **values of N data points** to  $a * x + b$


- Benefits from model-based storage:

- The storage needed for a model can be constant
- The structure of a time series is preserved intactly
- Queries can be answered directly from models

- Problems with model-based storage:

- The best model for a time series changes over time
- Long models for high compression increase latency

- Our contributions **remove both of these problems**



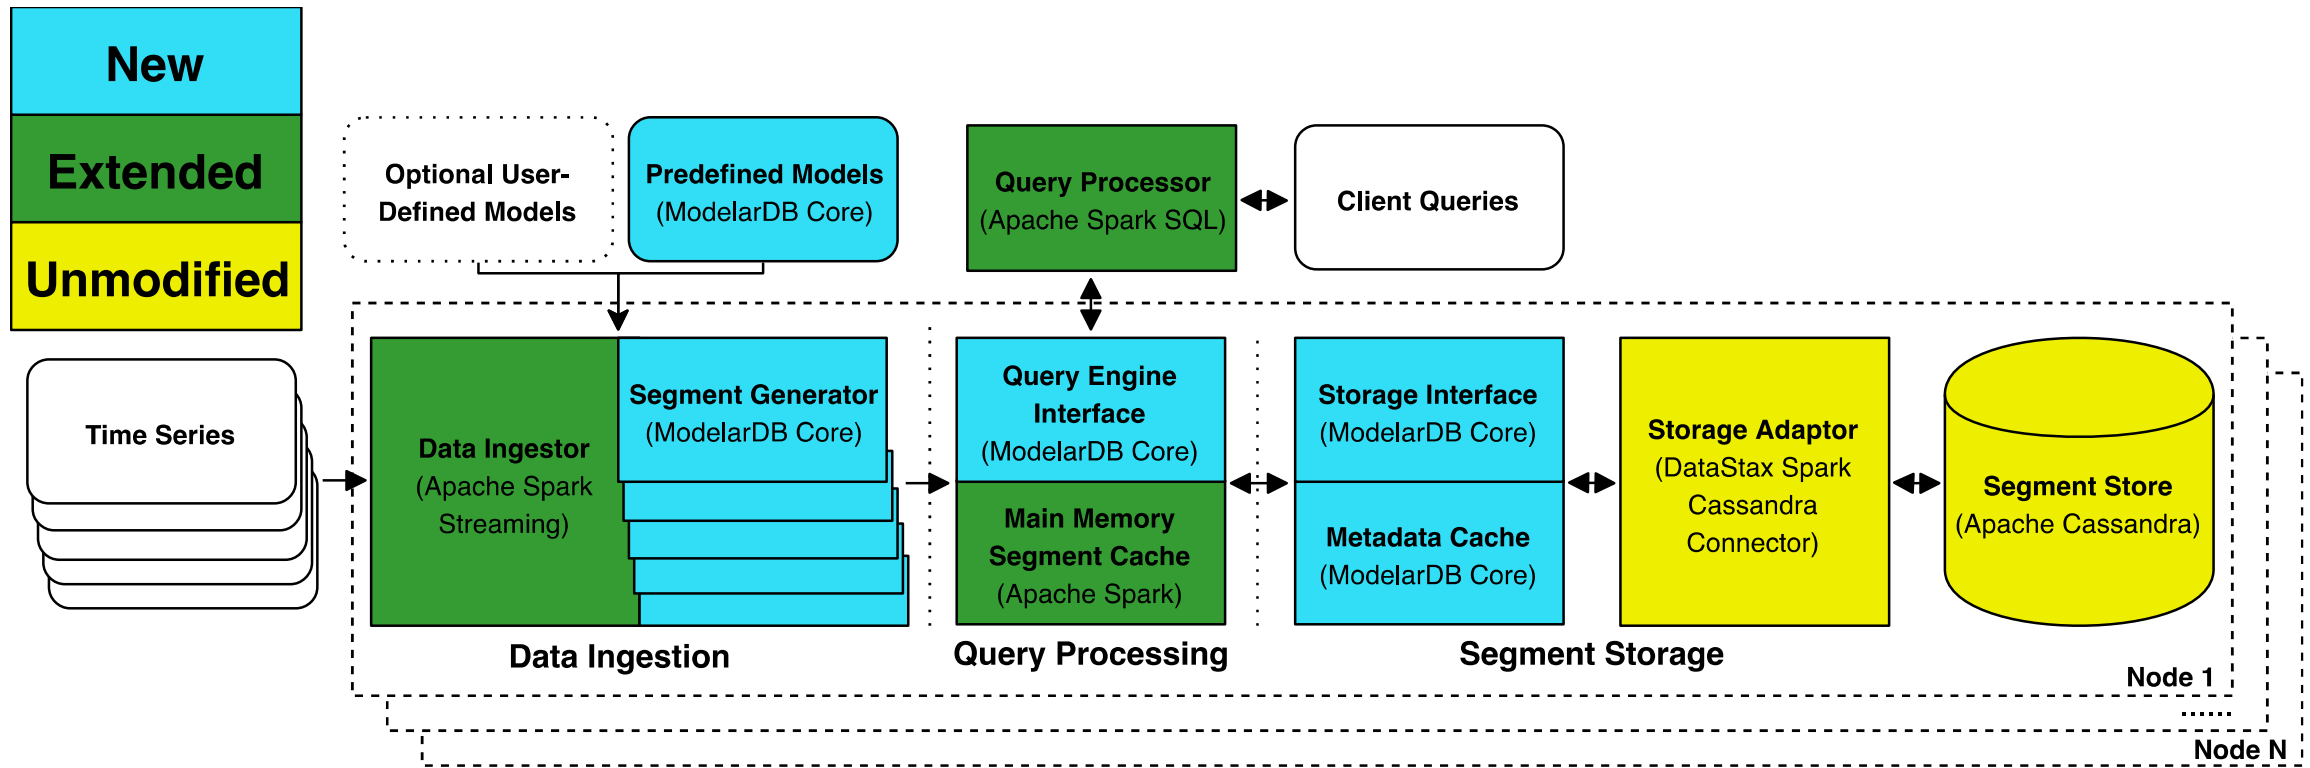
# Contributions

---



- A general-purpose architecture for a modular model-based time series management system (TSMS)
- An adaptive online algorithm for multi-model compression of time series
  - It is model-agnostic, extensible, allows gaps, and offers low latency and high compression
- A set of methods and optimizations for a model-based TSMS:
  - A database schema to store multiple time series as models
  - Methods to push-down predicates to a key-value store storing models
  - Methods to execute optimized aggregate functions directly on models
  - Use of static code-generation to optimize projections
  - Dynamic extensibility for adding models without recompiling the TSMS
- **ModelarDB** – an open-source implementation of our architecture
  - Available at [github.com/skejserjensen/ModelarDB](https://github.com/skejserjensen/ModelarDB) under version 2.0 of the Apache License

# Architecture




- All portable functionality is part of a separate library named **ModelarDB Core**
- Our implementation interfaces **ModelarDB Core** with **Spark** and **Cassandra**
  - ModelarDB can be deployed on unmodified instances of Spark and Cassandra

# Ingestion



- Models are incrementally fitted and emitted as part of **segments with metadata**:
  - Temporary Segment**: Holds an unfinished model cached in memory for low latency
  - Finalized Segment**: Holds a finished model cached in memory and persisted to disk
- Models are fitted in sequence until all would exceed the error bound:

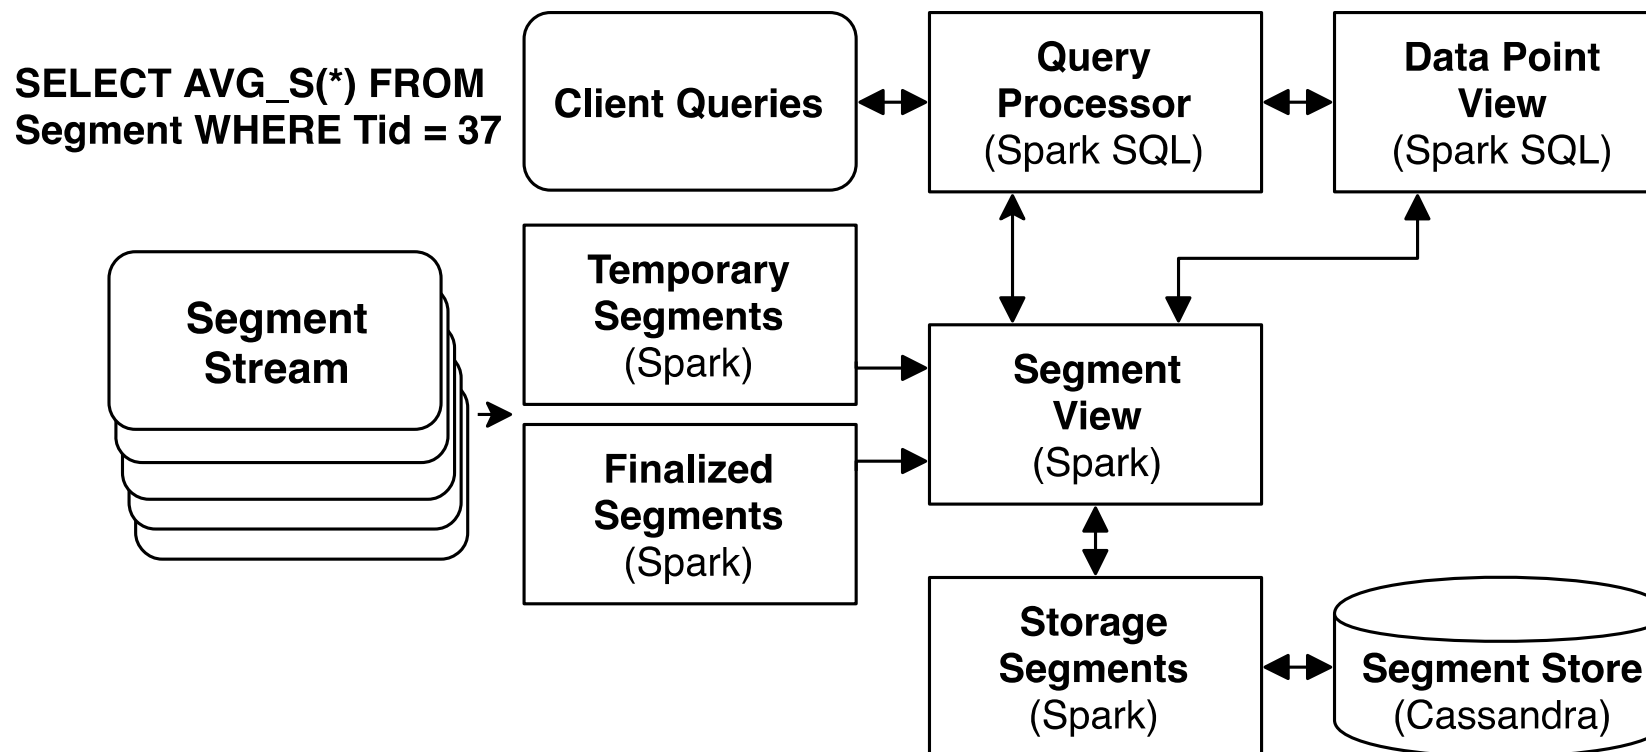


- ModelarDB Core includes four model types, users can **optionally add** more:
  - PMC-MR** (Constant), **Swing** (Linear), **Facebook** (Lossless), **Uncompressed** (None)

# Query Processing



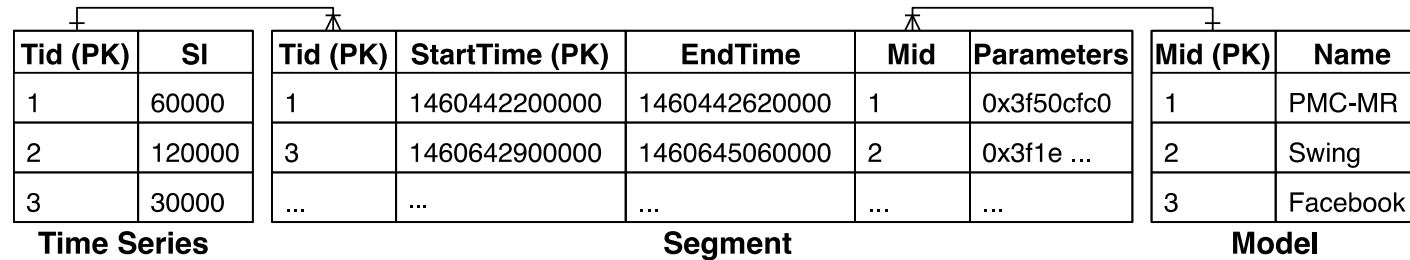
- ModelarDB uses **SQL for queries** to provide a familiar interface
- Two views are provided to allow for queries at different granularities:
  - **DataPoint View**: executes queries on data points reconstructed from segments
  - **Segment View**: efficiently executes aggregate queries directly on segments



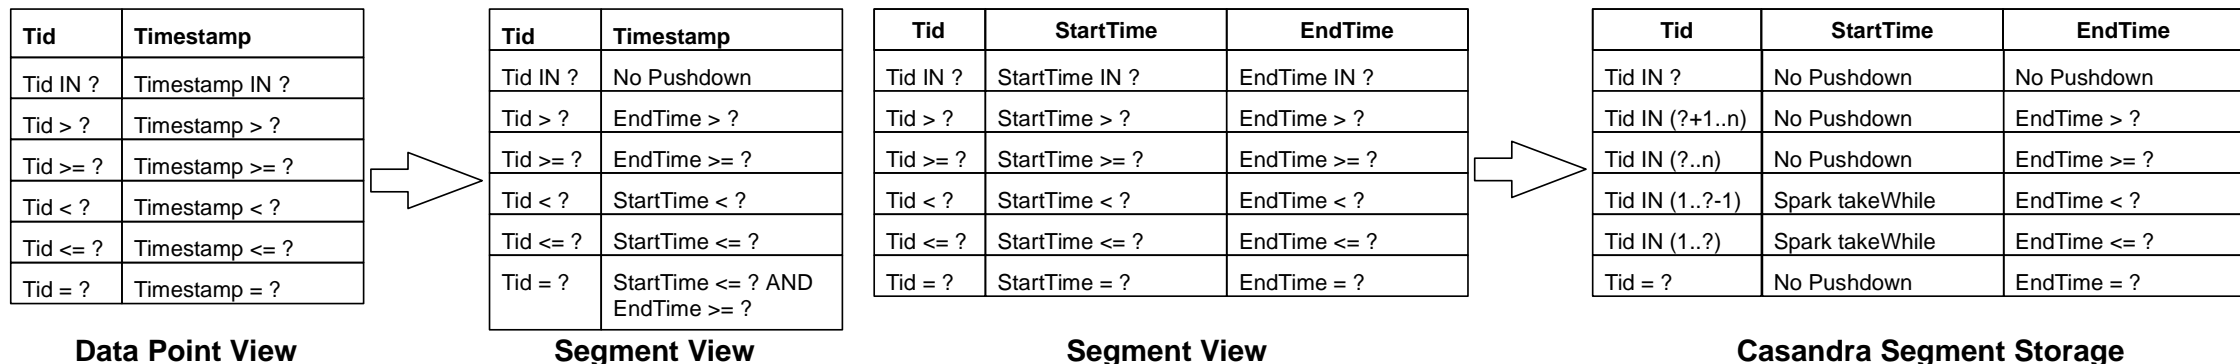
# Predicate Push-Down



- ModelarDB uses a three table schema for storing time series as segments



- ModelarDB performs predicate push-down as a multi-step procedure:
  - Data Point View:** Predicates are rewritten and pushed to the Segment View
  - Segment View:** Predicates are pushed without changes to the Storage Interface
  - Storage Interface:** Predicates are rewritten and pushed to the Segment Store
    - The Segment Store can have imprecise evaluation of the predicates (i.e., with false positives)





# Code Generation for Projection



- The overhead of projections are reduced using optimized lambda functions
- As the columns are known, the projection code is generated at compile time
- The correct function is found using a key created from the requested columns

```
1  def getDataPointGridFunction
2    (columns: Array[String]): (DataPoint => Row) = {
3    val target = getTarget(columns, dataPointView)
4    (target: @switch) match {
5      //Permutations of ('tid')
6      case 1 => (dp: DataPoint) => Row(dp.tid)
7      ...
8      //Permutations of ('tid', 'ts', 'value')
9      ...
10     case 321 => (dp: DataPoint) => Row(dp.value,
11       new Timestamp(dp.timestamp), dp.tid)
12   }
13 }
```

# Model-based aggregation

---



- Queries on the Segment View are executed directly on segments if possible
- Segments can implement optimized methods for aggregate queries
  - E.g., **sum for Swing** can be computed in **constant time** as shown below
- Aggregates are computed from reconstructed data points as a fallback

```
1 public double sum() {  
2     int timespan = this.endTime - this.startTime;  
3     int size = (timespan / this.SI) + 1;  
4     double first = this.a * this.startTime + this.b;  
5     double last = this.a * this.endTime + this.b;  
6     double average = (first + last) / 2;  
7     return average * size;  
8 }
```

# Evaluation - Storage



- 6 + 1 Laptops, **EH** (583 GiB, 100 ms), **ER** (488 GiB, 1 s), **EP** (339 GiB, 60 s)




Figure: Storage, EH




Figure: Storage, ER




Figure: Storage, EP

- ModelarDB provides better compression using model-based storage**
  - The best compression ratio is for high frequency data (EH, ER) and increases with error bound
  - Average error is 0.005% (EH), 2.5% (ER) and 0.73% (EP) for a 10% error bound
  - The paper shows ModelarDB degrades gracefully as outliers increase in number

# Evaluation - Adaptability

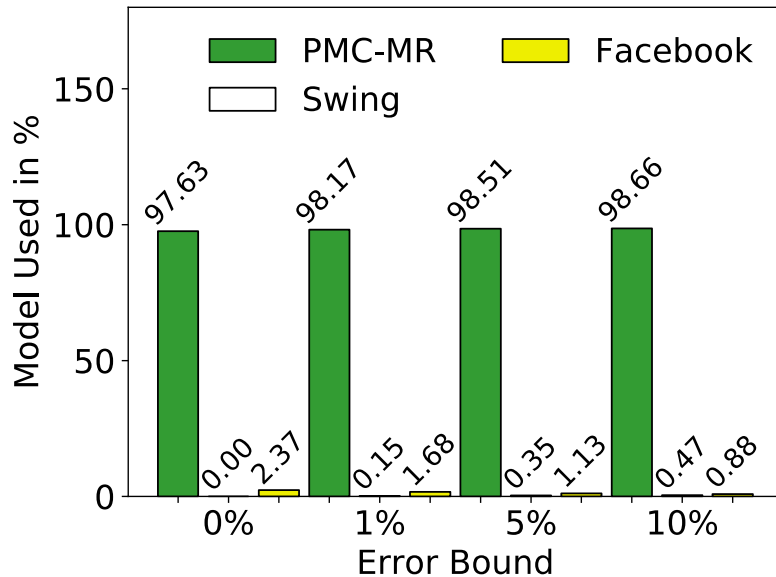


Figure: Adaptability, EH




Figure: Adaptability, ER






Figure: Adaptability, EP

- **ModelarDB chooses an appropriate model for each part of a series**
  - Different models are used for each data set and linear models are used with 0% error bound
  - The system is extensible and users can implement other models to increase adaptability


# Evaluation – Ingestion and Query Processing



**Figure:** Ingestion, ER



**Figure:** Aggregate Queries, ER



**Figure:** Scalability (Azure), ER

- **ModelarDB has fast ingestion, fast large aggregates and linear scalability**
  - Only InfluxDB, Cassandra, and ModelarDB can answer queries while ingesting data points
  - The paper shows ModelarDB is competitive with other systems for small scale queries

# Summary and Future Work

---



- **Summary:**

- Storing sensor data as **simple aggregates discards valuable information**
- **Model-based compression provides multiple benefits** over simple aggregates
- Proposed the model-based TSMS **ModelarDB** based on:
  - ◆ A general architecture for a modular model-based TSMS
  - ◆ An algorithm for online multi-model compression of time series
  - ◆ A set of methods and optimizations for a model-based TSMS
- Evaluation showed that **ModelarDB hits a sweet spot** by providing:
  - ◆ **Fast ingestion**
  - ◆ **Good compression**
  - ◆ **Fast, scalable online aggregate query processing**

- **Future Work:**

- Reduce storage required by storing correlated time series as one stream of segments
- Increase query performance through new methods that operate on time series as segments