

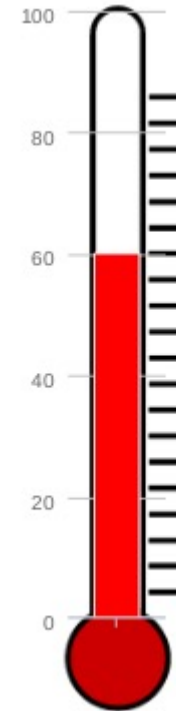
“Effektive metoder til at gemme og forespørge på store mængder tidsseriedata”

Christian Thomsen (chr@cs.aau.dk)

**Joint work with Søren Kejser Jensen and
Torben Bach Pedersen**

Center for Data-intensive Systems

Example – summer



Example – autumn



The challenge



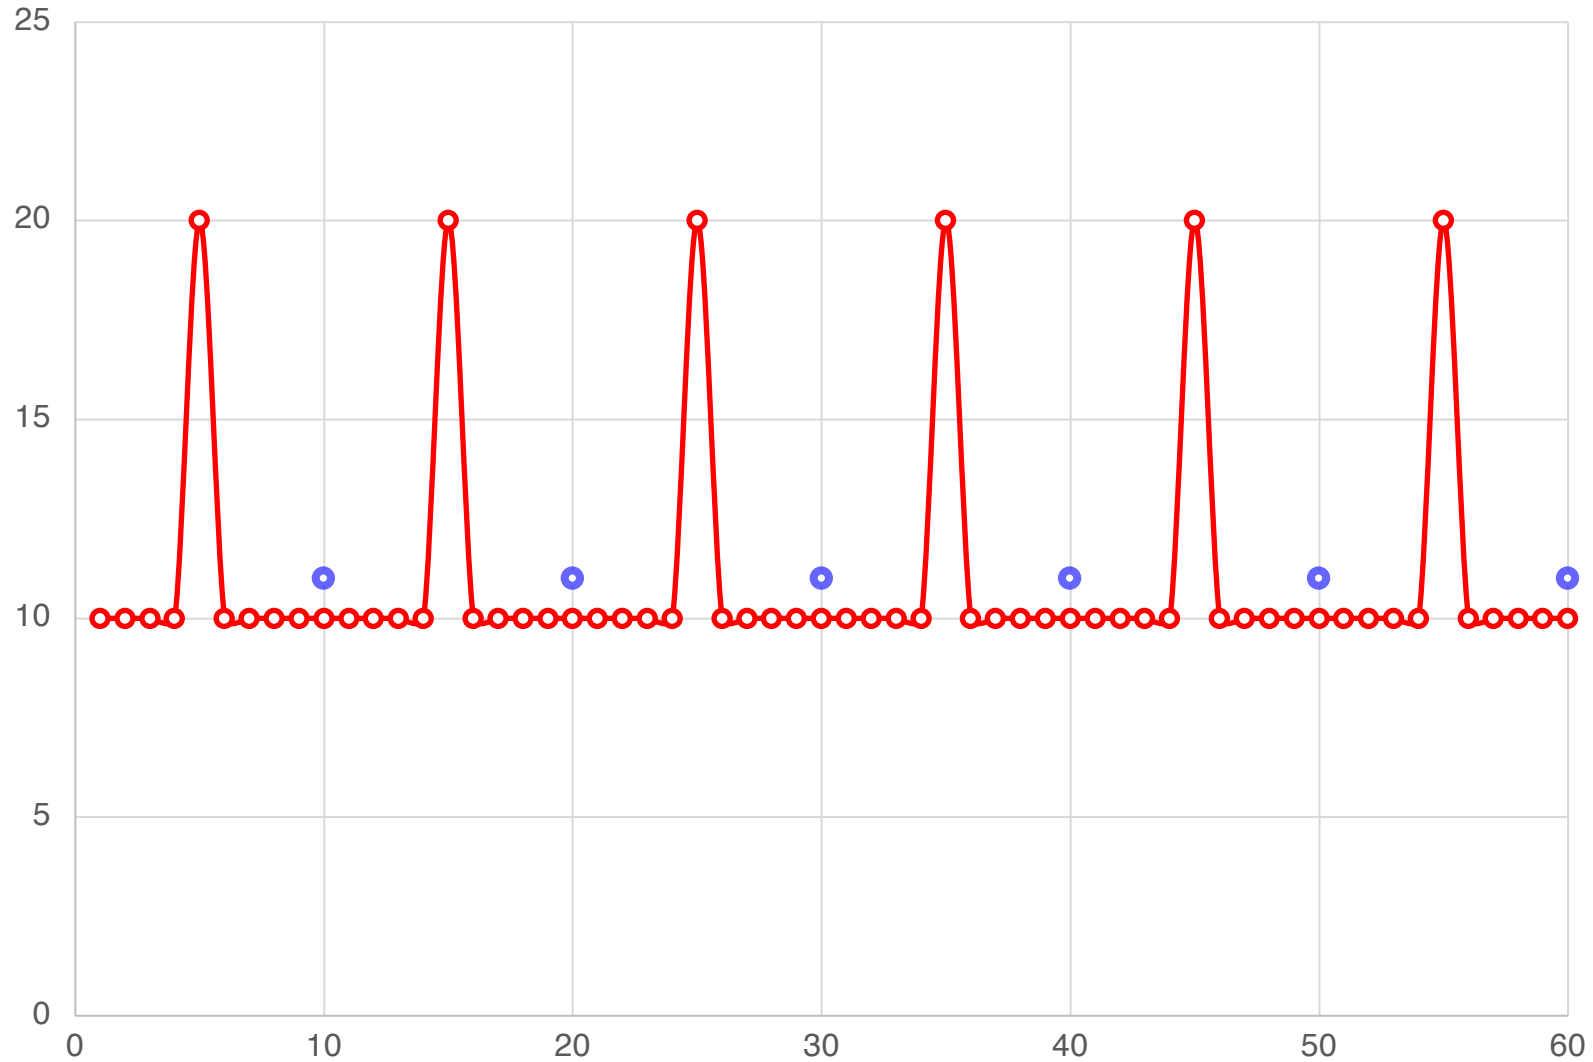
- Wind turbines and solar panels have a lot of sensors that can deliver data values several times per second
- A modern wind turbine has up to 6,500 streams
- This generates a lot of data
 - 10 reads/second, 4 bytes, 6,500 streams → ~20 GiB per day from one wind turbine
 - In practice, some sensors are sampled less often today, but a wind turbine still produces around 1GiB data per day
 - The sampling frequencies and amounts of data to store are increasing

The current situation



- The available information is currently not exploited or stored
- Many monitoring solutions consider few (~ 100) sensor streams and store only a single value for every x minutes (e.g., the average)
 - x is typically $\frac{1}{2}$, 1, 2, 5, or 10
- Important things might not be seen since outliers and fluctuations can be lost

Example of "missing the point" :-)



What we want to do...



- Store and use all available sensor data
- Support efficient aggregate queries on historical data
- Support analysis of data while it is being ingested
- Detect underperformance and other problems immediately
- Enable predictive maintenance

Why is that good? \$\$\$!



- For example, detect and fix a problem before the wind turbine breaks
- Reduced costs for service and spare parts
 - No over-time hours, crane booked in advance
- Service when there is little wind anyway
- Less downtime → more production
 - Delivery of a gearbox or wing can take months
- The service cost represents 11-30% of the onshore wind energy cost
- Global wind service revenue: 8 billion USD

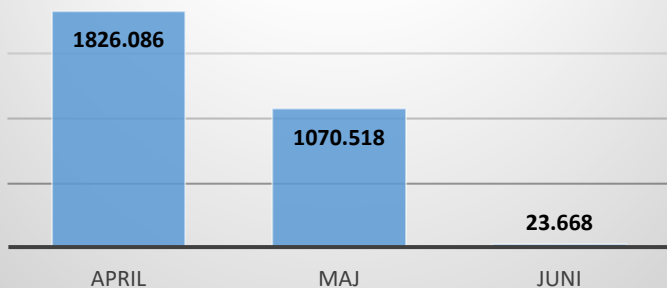
How we do it



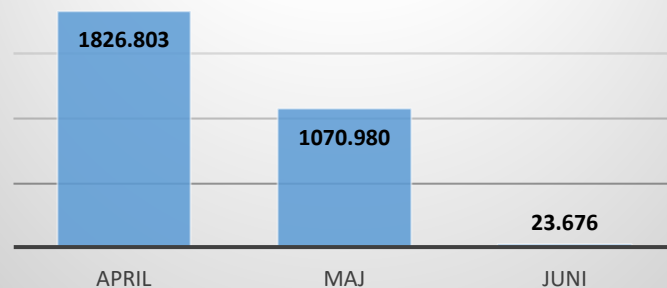
- Time-series can contain millions of points
- An efficient way to store and process them is to represent them by *models*
- We use a *model-based* approach for the time-series data
- A (user-defined) error-bound can be set
 - For example 5%, 1%, or even 0%
- Allowing an error in the representation can lead to better compression and performance

Ryd

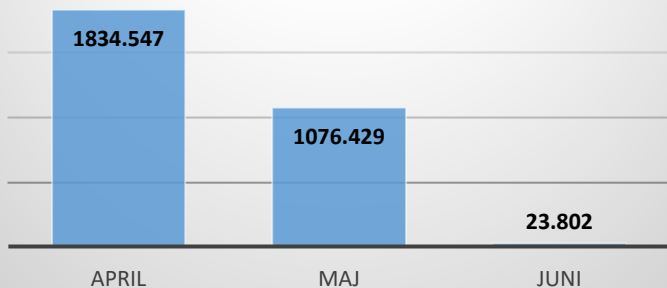
Ingen fejl - 0% (19:22)



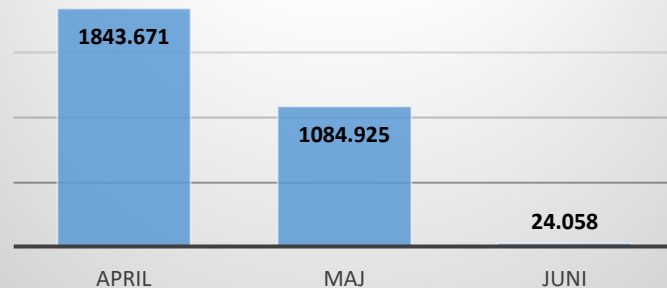
Op til 1% fejl tilladt (15:35)



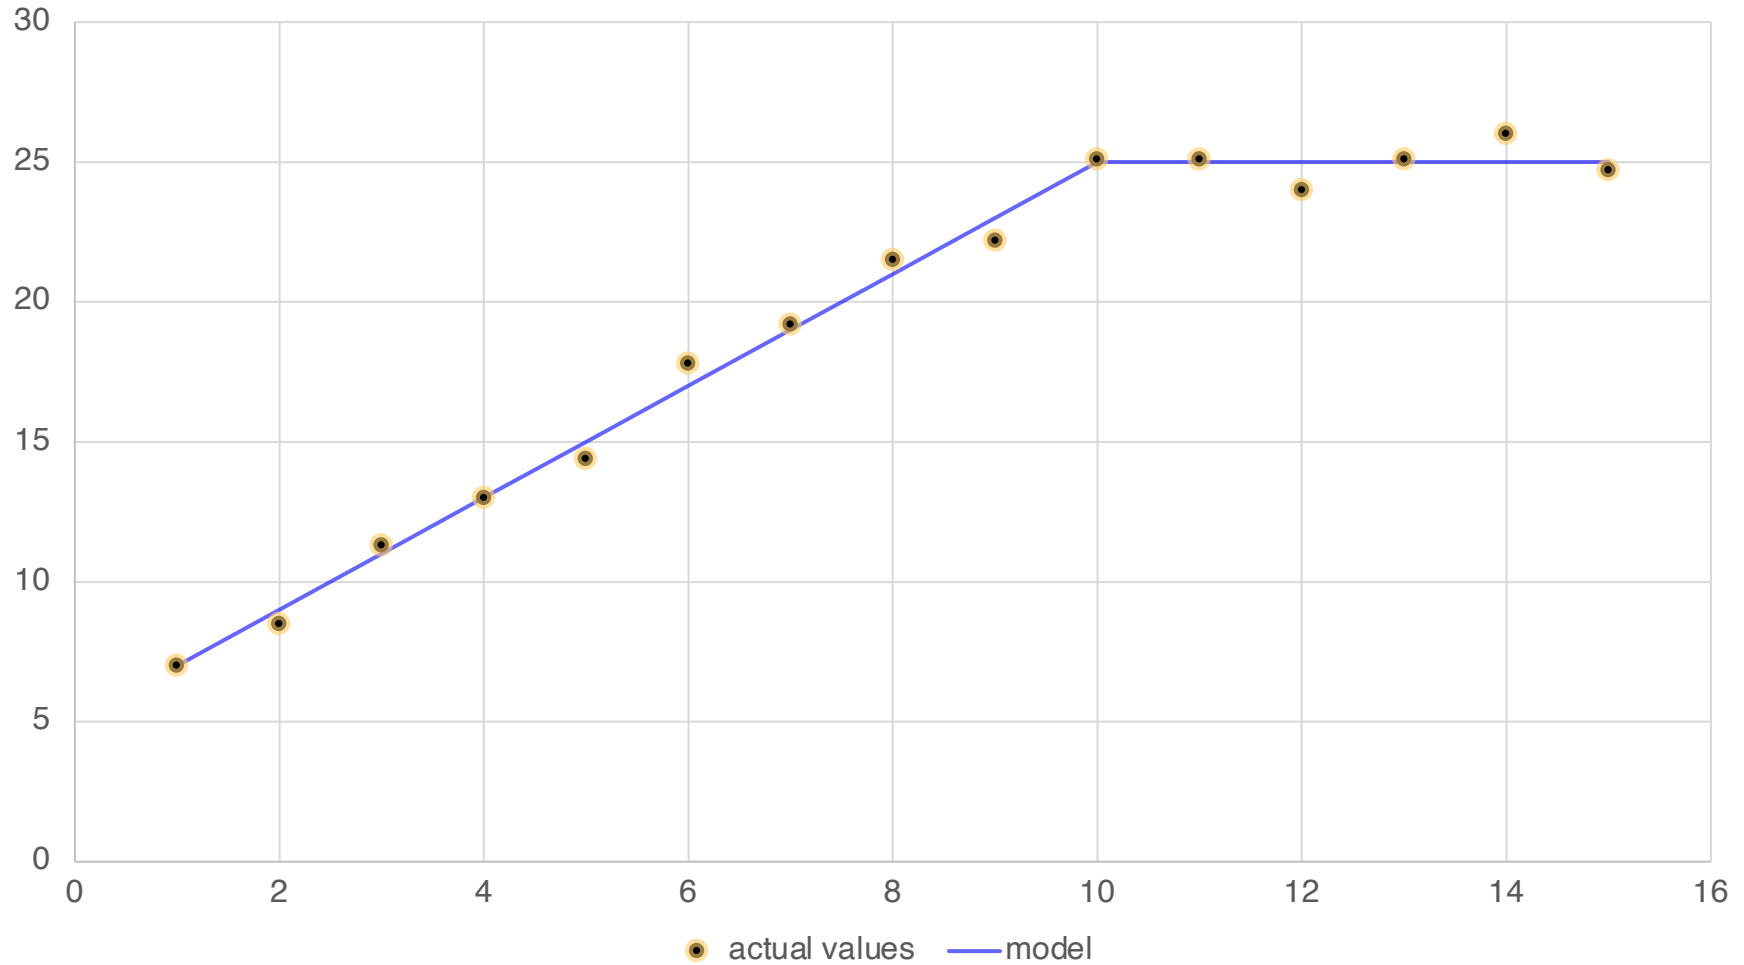
Op til 5% fejl tilladt (9:15)



Op til 10% fejl tilladt (7:38)



Simple example of models



ModelarDB



- We have developed the time series management system ***ModelarDB*** which uses models to store time series data
- Time series-specific functionality implemented in a system-agnostic library
- We have implemented some model types and the user can *optionally* add more
- ModelarDB adapts to the dataset and automatically picks the best model type to use for a given part of a time series
- Query processing and storage from existing systems
 - Apache Spark and Cassandra, respectively
 - Can be replaced by others

Storage requirements for a real-world data



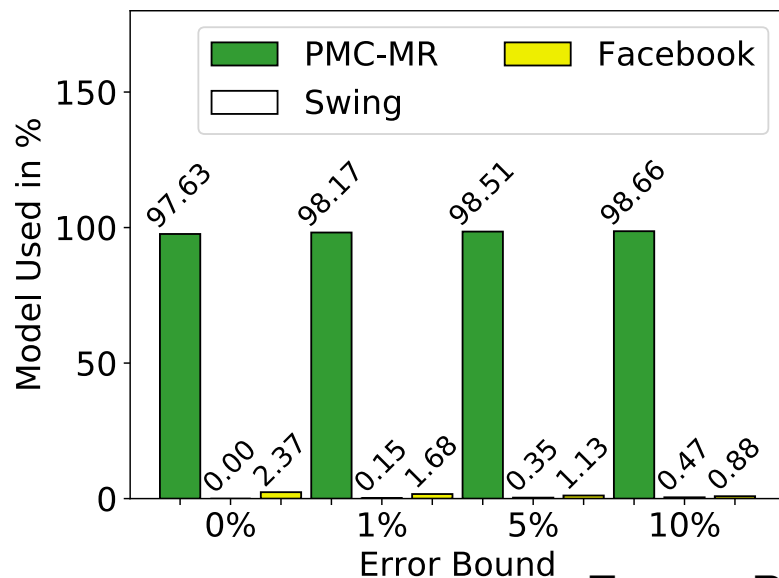
Storage Method	Size in GiB
CSV files	582.68
PostgreSQL 10.1	782.87
<i>RDBMS-X</i> (row)	367.89
<i>RDBMS-X</i> (column)	166.83
InfluxDB 1.4.2	4.33 – 4.44
Apache Parquet files	106.94
Apache ORC files	13.50
Apache Cassandra 3.9	111.89
ModelarDB	2.41 – 2.84

When the error bound is 10%, the actual average error is only 0.005% here!

Models used for different data sets



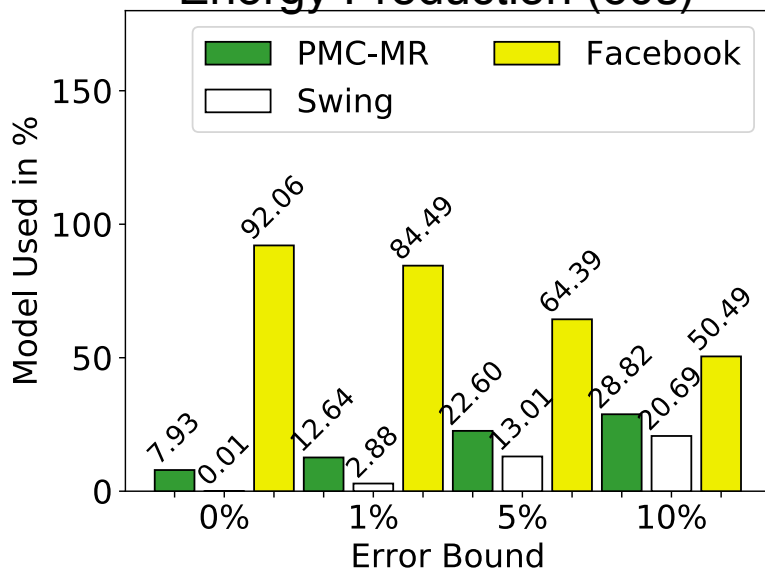
High-frequent production data (.1s)



Extended REDD data (1s)



Energy Production (60s)



Evaluation

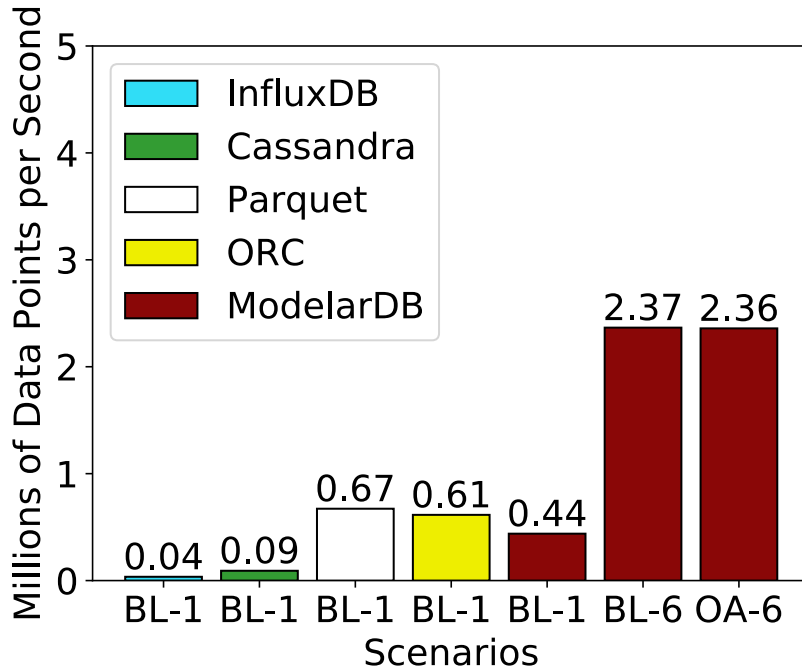


Figure: Ingestion, Ext. REDD

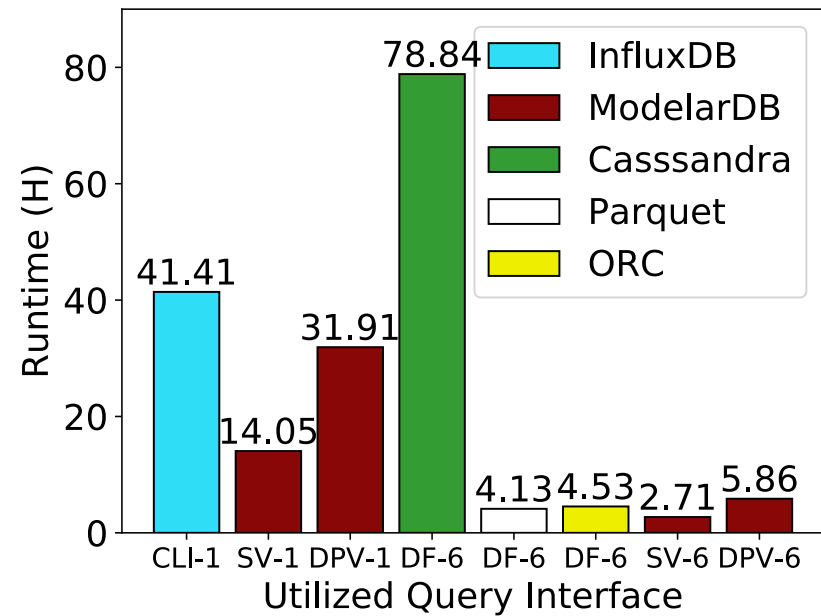


Figure: Aggregate Queries, Ext. REDD

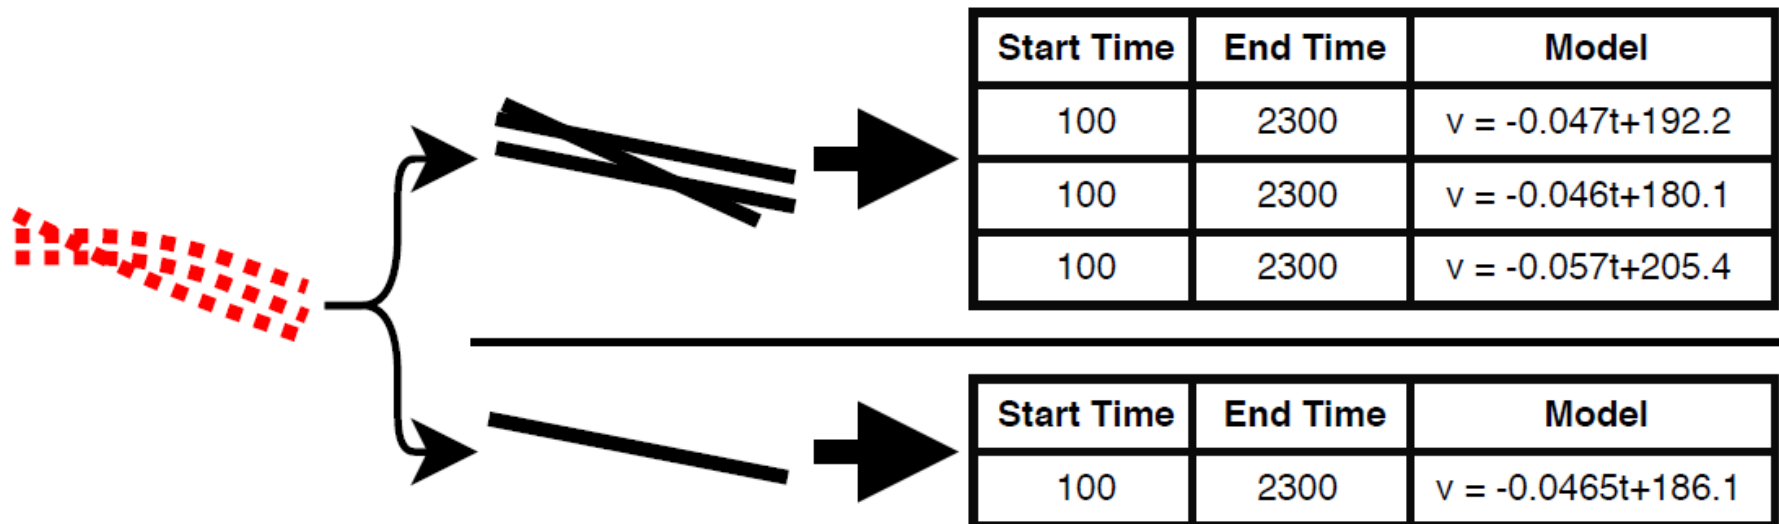
- Only InfluxDB, Cassandra, and ModelarDB can answer queries while ingesting data points

Performance summary



- ModelarDB provides support for fast ingestion, good compression, and fast large aggregate queries
- ModelarDB remains competitive for small aggregate and point/range queries
- Other systems are good for *one* of these, but not both
- ModelarDB also supports queries while ingesting data

Next step: Exploiting correlation

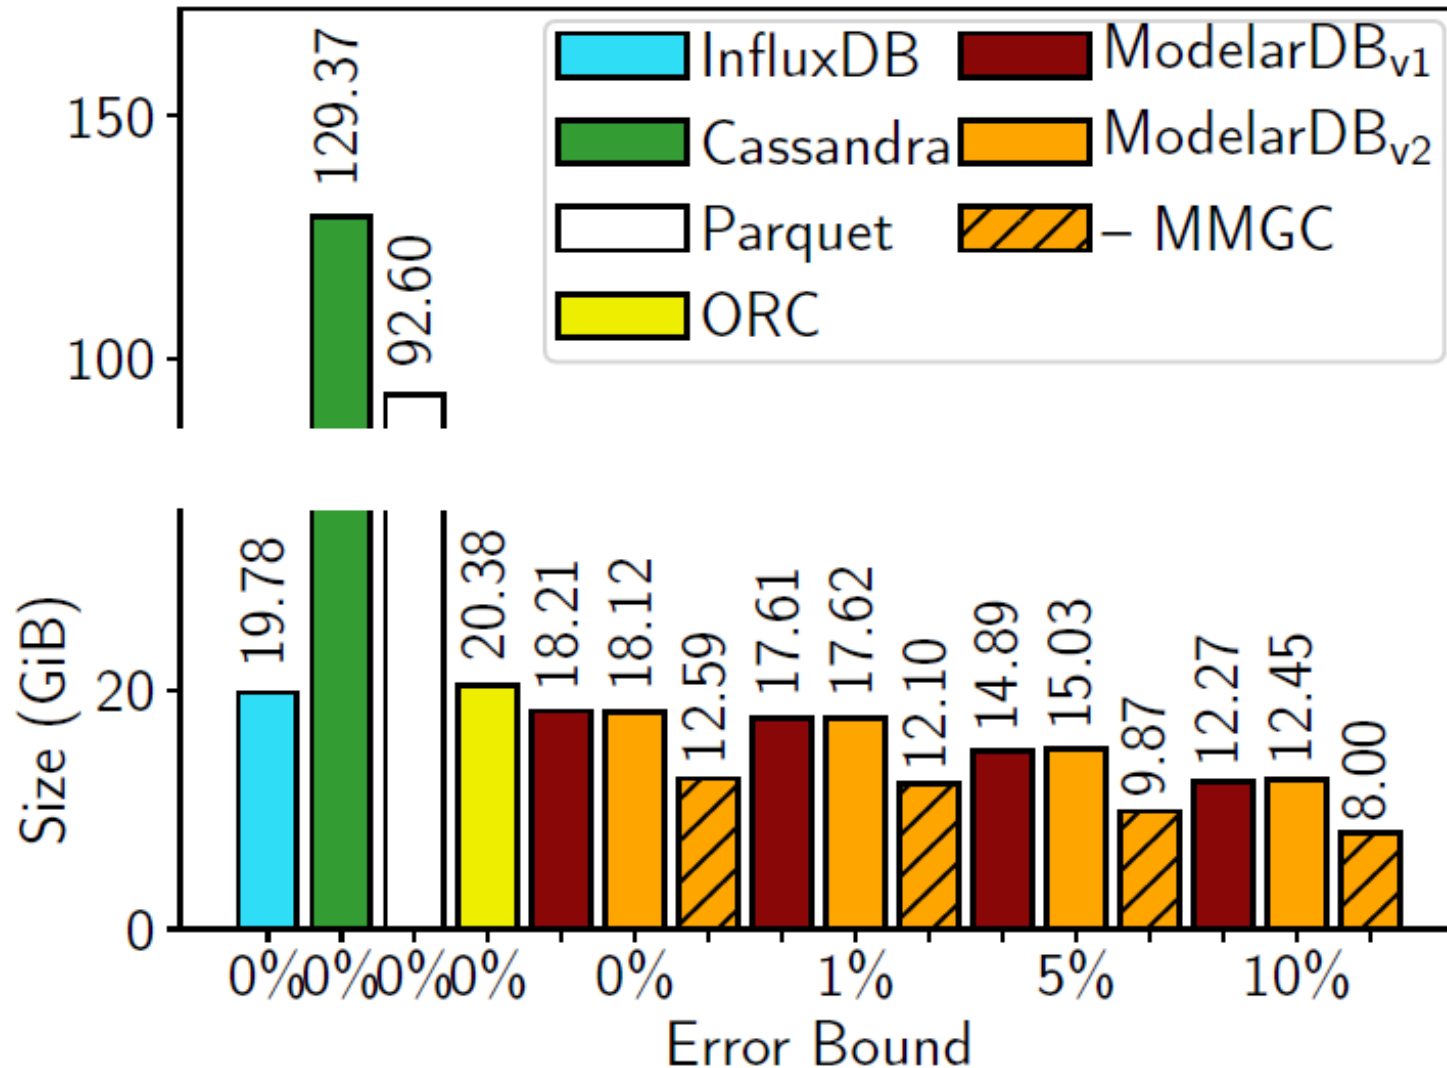


Specifying correlation



- Detecting correlation in data is an orthogonal problem
- We let the user hint correlation
- If the time series in a group cannot be represented by a single model, ModelarDB *splits* the group
 - Respects the error bound
- The time series can be *joined* again later

Evaluation



Conclusion and future work



- ModelarDB provides model-based compression within an error bound
- ModelarDB adapts to the dataset and compresses well by dynamically choosing among multiple models
- Good performance
- Integrated with Spark and Cassandra
- Future directions
 - Indexing to increase query performance further
 - Dynamic sampling
 - Advanced edge processing

More information



- S.K. Jensen, T.B. Pedersen, and C. Thomsen: "ModelarDB: Modular Model-Based Time Series Management with Spark and Cassandra", PVLDB 11(11), is available from <http://www.vldb.org/pvldb/vol11/p1688-jensen.pdf>
- S.K. Jensen, T.B. Pedersen, and C. Thomsen: "Scalable Model-Based Management of Correlated Dimensional Time Series in ModelarDB" is available from <https://arxiv.org/abs/1903.10269>

Acknowledgments



This research was supported by

- the DiCyPS center funded by Innovation Fund Denmark
- the GOFLEX project EU grant agreement No 731232
- and Microsoft Azure for Research