

# Model-Based Time Series Management at Scale

## PhD Thesis Defense

November 4<sup>th</sup>, 2019

Søren Kejser Jensen  
skj@cs.aau.dk

Center for Data Intensive Systems, Daisy  
Department of Computer Science  
Aalborg University  
Denmark

Supervisor: Torben Bach Pedersen  
Co-Supervisor: Christian Thomsen



# Agenda



Motivation and Publications

Time Series Management Systems

Model-Based Compression

Implementation of ModelarDB

Conclusion and Future Work



# Background

## Motivation and Publications

- Denmark targets being self-reliance on renewable energy by 2050.<sup>1</sup>
- 61.6% of Denmark's total energy production was renewable in 2016.<sup>2</sup>
- 71.8% of renewable energy produced in 2016 comes from Wind.<sup>2</sup>
- The energy produced by a wind turbine can be reduced due to, e.g.:
  - The wings of the wind turbine being covered in ice.
  - The coating on the wings deteriorating over time.
  - Dust from fields being harvested clogs the exhaust.
- Damaged parts can cause months of downtime as new are delivered.
- Through extensive monitoring, problems with wind turbines can be preemptively detected and corrected during scheduled maintenance.

---

<sup>1</sup>The Danish Energy Model - Innovative, efficient and sustainable

<sup>2</sup>Environmental report for Danish electricity and CHP for 2016 Status Year



# Background

## Motivation and Publications

### Meetings with manufacturers, owners, and energy traders:


- Modern turbines are monitored by up to 7,000 high-quality sensors.
- The sensors are installed with wired power and connectivity.
- Each sensor produces a data stream sampled to, e.g., a 10hz series.
- Collected measures include: Air Pressure, Humidity, Watt, Total Watt, Rotation Speed, Temperature, Wind Direction, Wind Speed.
- The time series are regular, cleaned, but missing values can occur.
- Aggregates with the following structure are the primary query type:
  - `SELECT {Aggregates} FROM {Table Name} WHERE  
time >= {Start} AND time < {End} {Optional Extra}  
GROUP BY time({Resolution}) {Optional Extra}`

# Challenges

## Motivation and Publications

### Meetings with manufacturers, owners, and energy traders:

- The storage needed makes storing the high-frequency data infeasible.
- Manufacturers of wind turbines gather petabytes of data each month.
- Simple aggregates (e.g. 10-minute averages) are stored instead of the high-frequency series, thereby removing fluctuations and outliers.



- Users believe problems can be found earlier with high-frequency data.
- Compression need only be lossless for some types of time series:
  - E.g., the amount of energy produced must be exact for billing.



# Publications in the Thesis

## Motivation and Publications

This thesis proposes methods and algorithms for model-based time series management at scale, all implemented in the open-source *ModelarDB*.

## Time Series Management Systems

[A] Time Series Management Systems: A Survey, **TKDE'17**

## Model-Based Time Series Management

[B] ModelarDB: Modular Model-Based Time Series Management with Spark and Cassandra, **PVLDB'18**

[C] Scalable Model-Based Management of Correlated Dimensional Time Series in ModelarDB, **Unpublished<sup>3</sup>**

[D] Demonstration of ModelarDB: Model-Based Management of Dimensional Time Series, **SIGMOD'19**

---

<sup>3</sup>Currently under submission



# Agenda

Motivation and Publications

Time Series Management Systems

Model-Based Compression

Implementation of ModelarDB

Conclusion and Future Work



# Paramount Properties

## Time Series Management Systems

Paramount properties for systems managing wind turbine data:


- **Distribution:** The system must be able to scale to many nodes.
- **Stream Processing:** Data points are arriving continuously as a regular time series and must be queryable with a short latency.
- **Compression:** High compression is needed for high-frequency data.
- **Efficient Retrieval:** Indexes or ordered storage for fast retrieval.
- **Approximate Query Processing:** Approximate answers can be accepted for some time series and enables use of lossy compression.
- **Extensibility:** Allows users with domain knowledge to implement new storage methods optimized specifically for their data sets.




# Architectures

## Time Series Management Systems


- Existing systems were surveyed with a focus on the six properties.
- The survey showed that three different architectures were common.



Internal Storage



External Storage



RDBMS Extensions



# Internal Storage

## Time Series Management Systems

*tsdb, FAQ, WearDrive, RINSE, Perera et al., Plato, Chronos, Pytsms, PhilDB*

### Advantages

- Each component can be optimized specifically for the system.
- Transfer of data between multiple sub-systems is not required.
- Simple to deploy as the system does not need external sub-systems.

### Disadvantages

- Cannot reuse external sub-systems already deployed.
- Long development times if embeddable sub-systems are not used.

### Missing Paramount Properties

- None of the systems can natively run distributed on a cluster.
- Most systems are not extensible so users cannot add new formats.

# External Storage

## Time Series Management Systems

*TSDS, SciDB, Respawn, SensorGrid, Guo et al., Tristan, Druid, Huang et al., Williams et al., Bolt, Storacle, Gorilla, Unnamed, servloTicy, BTrDB*

### Advantages

- External sub-systems can be reused to reduce development time, and simplify deployment if the sub-systems are already deployed.

### Disadvantages

- Transfer of data between multiple sub-systems is required.
- The query processor is restricted to the data store's interface.
- Complex to deploy due to the multiple separate sub-systems.

### Missing Paramount Properties

- Only very limited prototypes focus on general-purpose AQP.
- Most systems are not extensible so users cannot add new formats.



# RDBMS Extensions

## Time Series Management Systems

*TimeTravel, F<sup>2</sup>DB, Bakkalian et al.*

### Advantages

- Existing functionality can be reused to reduce development time.
- Transfer of data between multiple sub-systems is not required.
- Simple to deploy as the system does not need external sub-systems.

### Disadvantages

- Unused functionality like transactions add overhead and complexity.
- The extensions are restricted to the API provided by each RDBMS.

### Missing Paramount Properties

- None of the systems can natively run distributed on a cluster.
- TimeTravel and F<sup>2</sup>DB focus on forecasting, not compression.



# Agenda

Motivation and Publications

Time Series Management Systems

Model-Based Compression


Implementation of ModelarDB

Conclusion and Future Work

# Individual Time Series

## Model-Based Compression

- A time series can be stored efficiently as a sequence of models.
- A model is any representation that can reconstruct the sub-sequence of values it represents within a known error bound (possibly zero):
  - E.g.,  $v = a \times t + b$  can represent a sub-sequence using only  $a$  and  $b$ .
- Fitted models are stored with the necessary metadata in a segment.
- A model type finds a model's parameters by fitting it to data points.
- Different model types can represent different structures efficiently.




- A model type dictates the representation and error function used.

# Correlated Time Series

## Model-Based Compression

- A data set of time series can contain redundant information:
  - Co-located temperature sensors will produce similar values.
- In ModelarDB we consider time series correlated for a set of model types if compressing them as a group reduces the storage required.
- A list of model types fit models to data points, e.g., a constant (PMC-Mean), linear (Swing), and lossless (Gorilla) model type:



- The model providing the best compression ratio is written to disk.



# How to Group Time Series?

## Model-Based Compression

### Problems

- The model types used change the definition of correlation, so users should be able to create groups that match their model types.
- Resources required to compute correlation might not be available.
- A group of correlated series should be ingested together on a node.
- To prevent migration, groups should be made before ingestion starts.

### Solution

- Time series are grouped based on user hints given using primitives.
- The primitives can be combined and allow users to state that series are correlated based on their source or their dimensional hierarchy.
- Users can use their domain knowledge or analyze historical data.




# Static Grouping

## Model-Based Compression

### Grouping 9632 and 9634:

- From specific sources:  
9632 9634
- Sharing a specific member:  
Location 3 Aalborg
- Share members until a level:  
Location 3
- The dimensions distance:  
0.25




Location dimension for wind turbines

The most recent version of ModelarDB implements auto grouping based on the rule of thumb that using the shortest distance provides a benefit.

# Dynamic Grouping

## Model-Based Compression

- Correlated series might temporarily not produce similar values:
  - E.g., A temperature sensor covered by a shadow and one that is not.
- This can be efficiently detected as the compression ratio is reduced.
- Dynamically splitting and joining groups remedies this problem:



# Agenda



Motivation and Publications

Time Series Management Systems

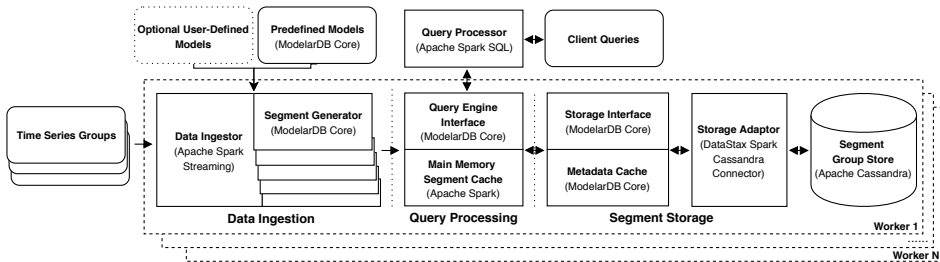
Model-Based Compression

Implementation of ModelarDB

Conclusion and Future Work

# ModelarDB Architecture

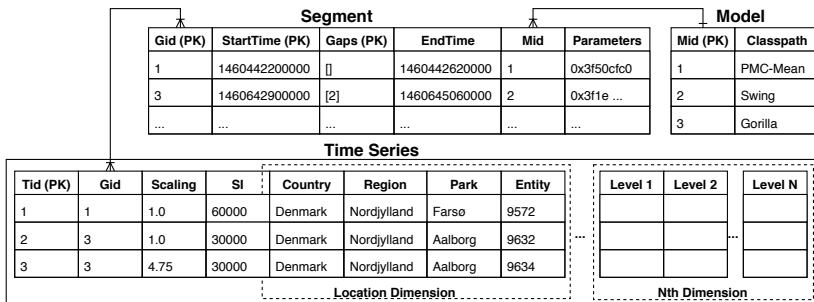
## Implementation of ModelarDB



- The general architecture consists of three sets of components: Data Ingestion, Query Processing, and Segment Storage.
- ModelarDB implements it using Apache Spark and Cassandra.
- External storage is used for extensibility and to simplify development.
- Extensible with regards to both data storage and query processing.

# Storage

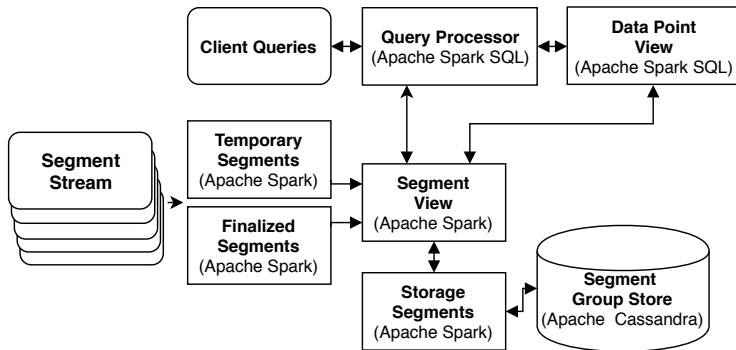
## Implementation of ModelarDB



- Time Series and Model store metadata for series and model types.
- Segment stores sub-sequences of series as segments with a model.

# Query Processing

## Implementation of ModelarDB



- Queries are given in SQL and executed on segments or data points.
- Predicate push-down and code-generation reduce execution time.
- For low latency, temporary segments are regularly cached in memory.

# Model and Segment API

## Implementation of ModelarDB

Model	
<code>new(Mid, Error, Limit)</code>	●
<code>append([Data_Point])</code>	●
<code>initialize([[Data_Point]])</code>	●
<code>parameters(Start_Time, End_Time, SI, [[Data_Point]])</code>	●
<code>get(Tid, Start_Time, End_Time, SI, Parameters, Offsets)</code>	●
<code>length()</code>	●
<code>size(Start_Time, End_Time, Resolution, [[Data_Points]])</code>	●
Segment	
<code>get(Timestamp, Index)</code>	●
<code>min()</code>	○
<code>max()</code>	○
<code>sum()</code>	○

- Methods marked ● are required while ○ methods are optional.
- Model and Segment are split to reduce the size of segment.
- The interfaces are used both internally and for user-defined types.



# Evaluation - Environment

## Implementation of ModelarDB

- The evaluation uses real-life data sets from the energy domain.
- Most experiments are performed on a modest local cluster.
- Scalability experiments are performed using Microsoft Azure.
- For comparison the evaluation includes the state-of-the-art big data systems and file formats most commonly used in industry.<sup>4</sup>
  - InfluxDB, Apache Cassandra, Apache Parquet, and Apache ORC
- No model-based time series management system is publicly available.
- Small, large, and multi-dimensional aggregates are evaluated using:
  - A command-line interface (CLI), a Spark DataFrame (S), the Segment View (SV), or the Data Point View (DPV).

---

<sup>4</sup>From meeting with companies, our survey, and <https://db-engines.com/en/>





# Evaluation - Data Sets


## Implementation of ModelarDB

- EP** 45,353 time series collected from energy producers with a sampling interval of 60s and occupying 339 GiB as gzipped CSV.
  - It contains two dimensions with each containing two levels:
    - **Production:** *Entity*  $\rightarrow$  *Type*  $\rightarrow$   $\top$
    - **Measure:** *Concrete*  $\rightarrow$  *Category*  $\rightarrow$   $\top$
  - Three time series with energy production are grouped per entity.
- EH** 286 time series collected from wind turbines with a sampling interval of 100ms and occupying 582.68 GiB as gzipped CSV.
  - It contains two dimensions with three and two levels respectively:
    - **Location:** *Entity*  $\rightarrow$  *Park*  $\rightarrow$  *Country*  $\rightarrow$   $\top$
    - **Measure:** *Concrete*  $\rightarrow$  *Category*  $\rightarrow$   $\top$
  - Time series with the same concrete measures are grouped per park.


# Evaluation - Ingestion and Storage

## Implementation of ModelarDB


- Results with groups of correlated time series are shown with stripes.



Ingestion Rate, EP



Storage Used, EP




Storage Used, EH


- Ingestion rate is 2.59–12.5 times faster than the industry formats.
- The storage required is reduced by 1.09–16.17 times for EP and by 1.30–64.63 times for EH compared to the industrial formats.
- Grouping series increases ingestion rate and can lower storage usage.

# Evaluation - Model Types

## Implementation of ModelarDB



Model Types, EP




Model Types, EH


- ModelarDB automatically adapts to each data set and error bound.
- Grouping increases ModelarDB's use of the lossless model type.

# Evaluation - Aggregate Queries


## Implementation of ModelarDB



Large Scale, EP



Small Scale, EP




Small Scale, EH

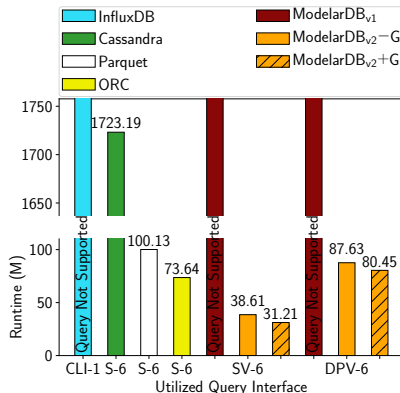
- Grouping time series decreases query time for large scale aggregate queries, but increases query time for small scale aggregate queries.
- Experiments run on Azure show that ModelarDB scales linearly.

# Evaluation - Multi-Dimensional Queries EP

## Implementation of ModelarDB



Month and Category, EP




Month and Concrete, EP


- Grouping reduces query time as each query only reads whole groups.

# Evaluation - Multi-Dimensional Queries EH

## Implementation of ModelarDB



Month and Park, EH



Month and Entity, EH

- Grouping increases query time as the cluster is not fully utilized.

# Agenda



Motivation and Publications

Time Series Management Systems

Model-Based Compression

Implementation of ModelarDB

Conclusion and Future Work



# Summary of Contributions

## Conclusion and Future Work

- A survey of existing time series management systems.
- An architecture for a model-based time series management systems.
- A model-agnostic compression algorithm for individual time series.
- A general schema for storing multiple time series as models.
- Methods for executing aggregate queries on user-defined models.
- Optimizations using predicate push-down and code-generation.
- Compression of correlated time series using multiple model types.
- Extensions of our methods from individual to correlated time series.
- Methods for partitioning time series into groups of correlated series.
- Methods for executing multi-dimensional aggregates on models.
- The extensible and modular implementation of ModelarDB.





# Conclusion

## Conclusion and Future Work

- Proposed a model-based approach for management of time series motivated by need for more detailed monitoring of wind turbines.
- Our ModelarDB system fulfills the paramount requirements:
  - **Distribution:** The system scales linearly to at least 32 nodes.
  - **Stream Processing:** Data is ingested and can be queried online.
  - **Compression:** Models efficiently represent each sub-sequence.
  - **Efficient Retrieval:** Data is partitioned and ordered for fast retrieval.
  - **AQP:** Allowing approximate answers reduce query response time.
  - **Extensibility:** Model types, query processor, and storage can change.
- ModelarDB hits a sweet spot and offers very fast ingestion, good compression, and fast, scalable online aggregate query processing.
- The contributions of this thesis increases the scale at which large time series can be collected, stored, and analyzed.



# Future Work

## Conclusion and Future Work

### Support new domains:

- Support time series with irregular sampling intervals.
- Evaluate our methods with data series from other domains.

### Ingestion and storage optimizations:

- Support ingestion directly at the sources to save bandwidth.
- Design a native storage format optimized for model-based storage.
- Inform the user if the current model types provide poor compression.

### Query processing:

- Index model values to answer queries faster and enable high level analytics directly on models.

### Automated parameter selection:

- Reduce the number of parameters to only an error bound.

# Our Modest Cluster

