

Extreme-Scale Model-Based Time Series Management with ModelarDB

Professor Torben Bach Pedersen
Aalborg University and ModelarData
tbp@cs.aau.dk

Joint work with Søren Kejser Jensen and Christian Thomsen

Center for Data-intensive Systems

A Tale of Some Really Big and Fast Data...




- Denmark is **world no. 1** i wind energy
 - Come and visit and you will feel why 😊
 - No 1 turbine maker Vestas is DK, No 2 Siemens Gamesa has most R&D in DK
 - World record electricity from wind: **50% 2019**, going towards 100% wind/solar/..
- Wind turbines
 - **500 sensors** -> more than **2500 derived data streams**
 - 8 byte values sampled at 100+ Hz, 100+ turbines in a park
 - $100 * 100 * 2500 = 25$ million values/second = **200+ MB/sec**
 - $200 \text{ MB} * 3600 * 24 = 17.5+$ TB/day = **8+ PB/year/park**
 - They want to store **20+ years for 1000s of parks...**
- Industry state of the art: **500 col SQL tables with 10 min avg...**
 - Makes high-frequency series impossible - how can we improve?
- Data characteristics:
 - **Regular sampling interval, out-of-order corrected, short gaps**

Observations




- From meetings with manufacturers, owners, and energy traders:
 - Turbines have **high-quality sensors** with **wired power and connectivity**
 - The storage needed makes **storing high-frequency sensor data infeasible**
 - Simple aggregates (e.g. **10-minute averages**) are **stored instead** of the **high-frequent series**, thereby **removing useful fluctuations and outliers**
- Many of the collected time series are **correlated** with each other
- They can be stored within a **user-defined error bound (possibly 0%)**
- **Metadata** is also stored and **aggregates** are the **primary query type**



The current situation




- The available information is currently not exploited or stored
- Many monitoring solutions consider few (~ 100) sensor streams and store only a single value for every x minutes (e.g., the average)
 - x is typically $\frac{1}{2}$, 1, 2, 5, or 10



- Important things might not be seen since **outliers and fluctuations are lost**

Example of "missing the point" :-)



What we want to do...



- Store and use **all available sensor data**
- Support **efficient aggregate queries** on historical data
- Support **analysis** of data **while** it is being **ingested**
- **Detect** underperformance and other **problems immediately**
- Enable **predictive maintenance**

Why is that good? € + CO2




- For example, detect and **fix a problem before** the wind **turbine breaks**
- **Reduced costs** for service and spare parts
 - No over-time hours, crane booked in advance
- Service when there is little wind anyway
- Less downtime → more production
 - Delivery of a gearbox or wing can take months
- The **service cost** represents **11-30% of** the **onshore wind energy cost**
- Global wind service revenue: **8 billion USD**
- **Most** importantly: making wind power cheaper **helps the Green Transition**
 - **Onshore wind** is already the **cheapest way** to install **new generation capacity**

How we do it



- Time-series can contain millions of points
- An efficient way to store and process them is to represent them by ***models***
- We use a ***model-based*** approach for the time-series data
- A **(user-defined) error-bound** can be set
 - For example **5%, 1%, or even 0%**
- **Allowing an error** in the representation can lead **to better compression** and **performance**

Simple example of models



More observations and some first results



- Wind turbines produce big amounts of **high-quality sensor data**
- Data is collected as **regular time series** with **only few gaps** without values
- **High frequency** could benefit analysis but:
 - High frequency data **cannot be ingested fast enough**
 - High **query processing time limits use of historical data**
 - **Unfeasible** high **amounts of storage** are required
- Storage of real-life wind turbine data:

Storage Method	Size (GiB)	Storage Method	Size (GiB)
PostgreSQL 10.1	782.87	CSV Files	582.68
RDBMS-X - Row	367.89	Apache Parquet Files	106.94
RDBMS-X - Column	166.83	Apache ORC Files	13.50
InfluxDB 1.4.2 - Tags	4.33	Apache Cassandra 3.9	111.89
InfluxDB 1.4.2 - Measurements	4.33	<i>ModelarDB</i>	2.41 - 2.84

Model-based storage of time series



- A **model** is a **lossy** or **lossless** representation of a time series
- E.g., a linear function reduces the **values of N data points** to $a * x + b$


- Benefits from model-based storage:

- The **storage** needed for a model can be **constant**
- The **structure** of a time series is **preserved** intactly
- **Queries** can be **answered directly from models**

- Problems with model-based storage:

- The **best model** for a time series **changes over time**
- **Long models** for high compression **increase latency**

- Our contributions **remove both of these problems**

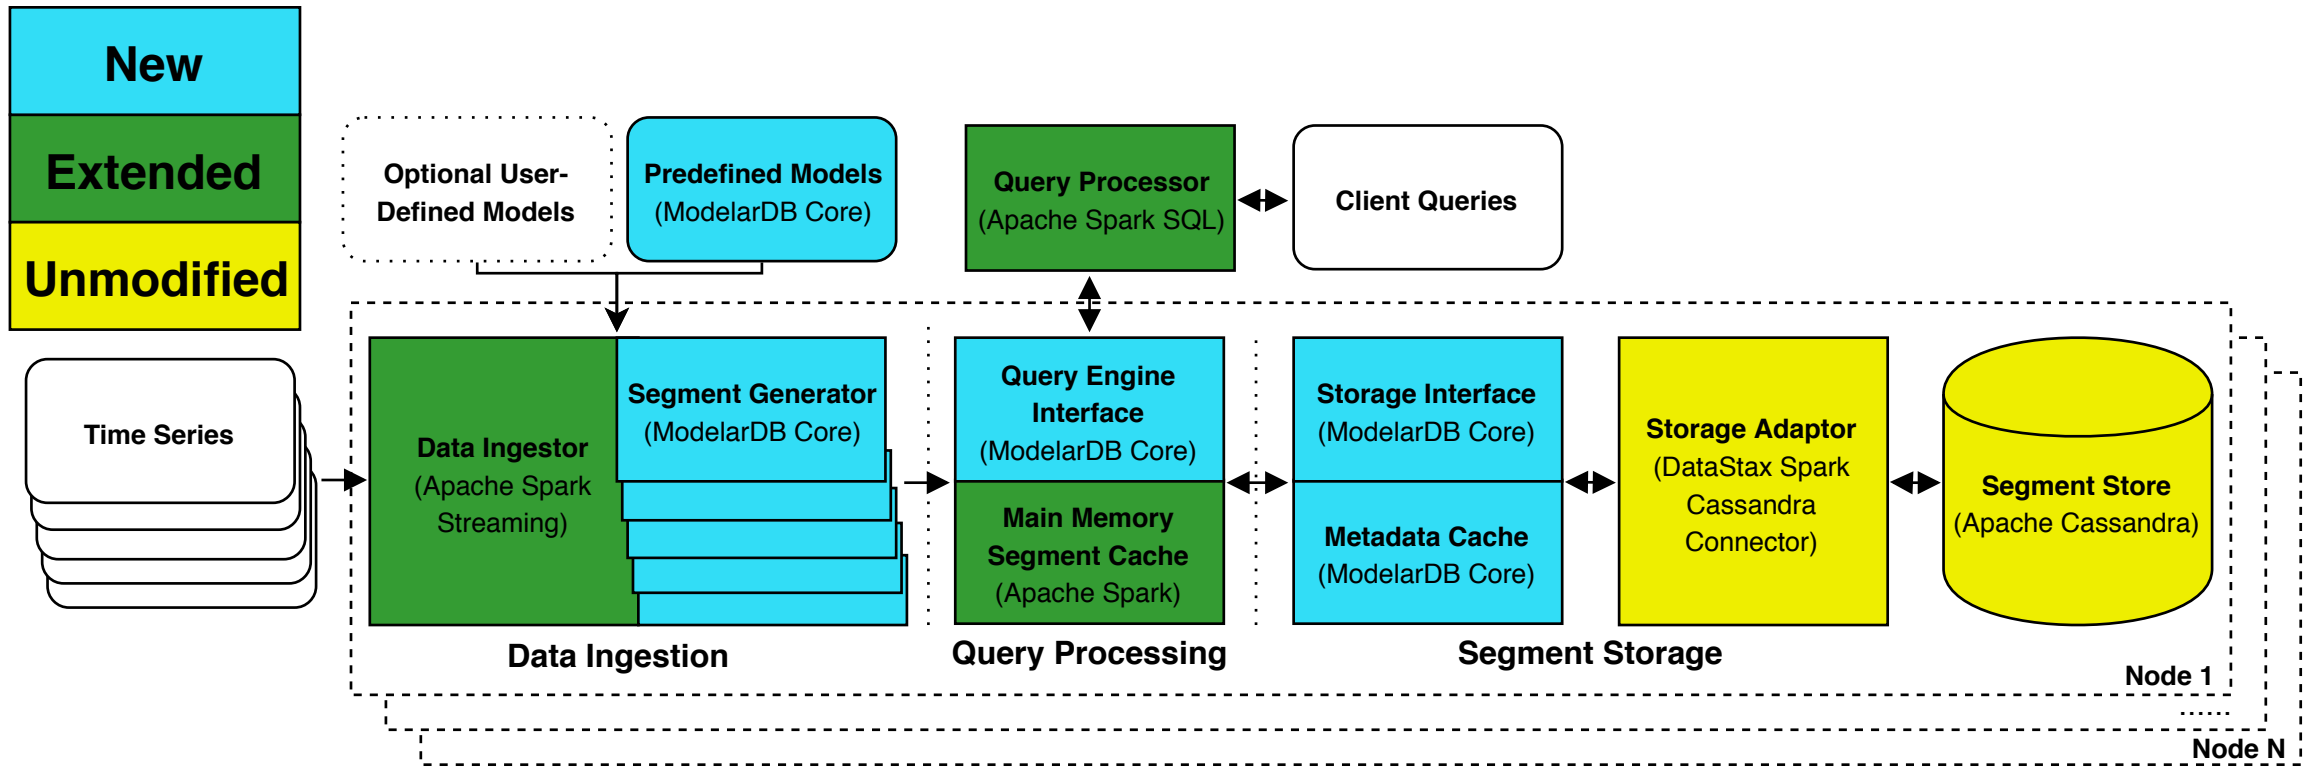


ModelarDB Contributions



- A **general-purpose architecture** for a modular model-based time series management system (TSMS)
- An **adaptive online algorithm for multi-model compression** of time series
 - **Model-agnostic, extensible, allows gaps + offers low latency and high compression**
- A set of **methods and optimizations** for a model-based TSMS:
 - A **database schema** to store multiple time series as models
 - Methods to **push-down predicates** to a key-value store storing models
 - Methods to **execute optimized aggregate functions** directly on models
 - Use of static code-generation to **optimize projections**
 - **Dynamic extensibility** for **adding models** without recompiling the TSMS
- **ModelarDB** – an **open-source** implementation of our architecture
 - Available at github.com/skejserjensen/ModelarDB under version 2.0 of the Apache License

Architecture




- All portable functionality is part of a separate library named **ModelarDB Core**
- Our implementation interfaces **ModelarDB Core** with **Spark** and **Cassandra**
 - ModelarDB can be deployed on unmodified instances of Spark and Cassandra

Ingestion



- Models are incrementally fitted and emitted as part of **segments with metadata**:
 - Temporary Segment**: Holds an unfinished model cached in memory for low latency
 - Finalized Segment**: Holds a finished model cached in memory and persisted to disk
- Models are fitted in sequence until all would exceed the error bound:




- ModelarDB Core includes four model types, users can **optionally add** more:
 - PMC-MR** (Constant), **Swing** (Linear), **Facebook** (Lossless), **Uncompressed** (None)

More details on ingestion



- **Fitting** seems as a **black box** to **support user-defined models**
 - Each **model** implements representing data points+measure error
 - **Model** is **passed data points** while user-def **error-bound** holds
- Models provide **trade-off between compression and latency**
 - **Longer models** give **better compression** but **higher latency**
- Example: **max latency of 3 data points (ye)**
 - Single model (**linear**) used to represent data points
 - Multiple models: the next model is evaluated when this fails



Query Processing Example




- Example using the Segment View

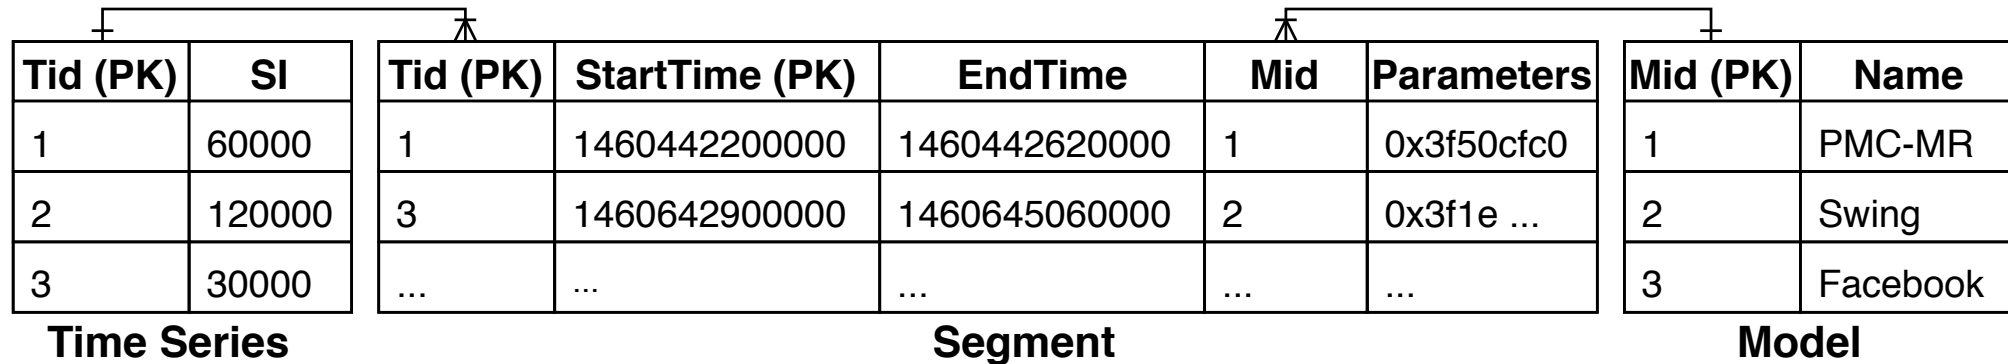
- Results is **SF₃** and **ST₆**
- SF₃** resides on **disk**
- ST₆** resides **in memory**

- Abbreviations:

- RS**: Result Set
- ST**: Temporary Segment
- SF**: Finalized Segment



Storage




- **Time Series**: Stores time series metadata
- **Model**: Stores model types utilized for segments
- **Segment**: Stores segments emitted for each time series
- The bulk of the data is stored as part of the segment table

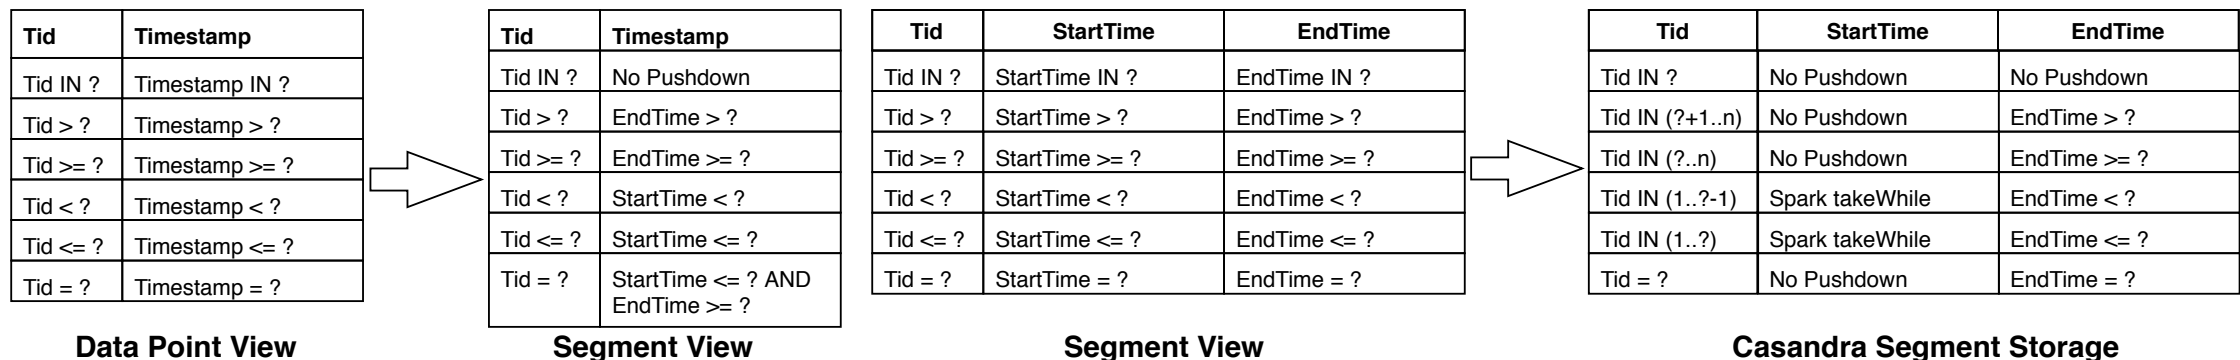
Predicate Push-Down



- ModelarDB uses a three table schema for storing time series as segments



- ModelarDB performs predicate push-down as a multi-step procedure:
 - Data Point View:** Predicates are rewritten and pushed to the Segment View
 - Segment View:** Predicates are pushed without changes to the Storage Interface
 - Storage Interface:** Predicates are rewritten and pushed to the Segment Store
 - The Segment Store can have imprecise evaluation of the predicates (i.e., with false positives)



Code Generation for Projection



- **Overhead of projections are reduced** using **optimized lambda functions**
- As the **columns are known**, the projection **code is generated at compile time**
- The correct function is found using a key created from the requested columns

```
1  def getDataPointGridFunction
2    (columns: Array[String]): (DataPoint => Row) = {
3    val target = getTarget(columns, dataPointView)
4    (target: @switch) match {
5      //Permutations of ('tid')
6      case 1 => (dp: DataPoint) => Row(dp.tid)
7      ...
8      //Permutations of ('tid', 'ts', 'value')
9      ...
10     case 321 => (dp: DataPoint) => Row(dp.value,
11       new Timestamp(dp.timestamp), dp.tid)
12   }
13 }
```

Model-based aggregation



- **Queries on Segment View** executed **directly on segments** if possible
- Segments can **implement optimized methods for aggregate queries**
 - E.g., **sum for Swing** can be computed in **constant time** as shown below
- Aggregates are computed from reconstructed data points as a **fallback**

```
1 public double sum() {  
2     int timespan = this.endTime - this.startTime;  
3     int size = (timespan / this.SI) + 1;  
4     double first = this.a * this.startTime + this.b;  
5     double last = this.a * this.endTime + this.b;  
6     double average = (first + last) / 2;  
7     return average * size;  
8 }
```

Extensibility with user-defined models



Table 2: Interface for models and segments, ● is a required method and ○ is an optional method

Model

<code>new(Error, Limit)</code>	●	Return a new model with the user-defined error bound and length limit.
<code>append(Data Point)</code>	●	Append a data point if it and all previous do not exceed the error bound.
<code>initialize([Data Point])</code>	●	Clear the existing data points from the model and append the data points from the list until one exceeds the error bound or length limit.
<code>get(Tid, Start Time, End Time, SI, Parameters, [Gap])</code>	●	Create a segment represented by the model from serialized parameters.
<code>get(Tid, Start Time, End Time, SI, [Data Point], [Gap])</code>	●	Create a segment from the models state and the list of data points.
<code>length()</code>	●	Return the number of data points the model currently represents.
<code>size()</code>	●	Return the size in bytes currently required for the models parameters.

Segment

<code>get(Timestamp, Index)</code>	●	Return the value from the underlying model that matches the timestamp and index, both are provided to simplify implementation of this interface.
<code>parameters()</code>	●	Return the segment specific parameters necessary to reconstruct it.
<code>sum()</code>	○	Compute the sum of the values of data points represented by the segment
<code>min()</code>	○	Compute the minimum value of data points represented by the segment.
<code>max()</code>	○	Compute the maximum value of data points represented by the segment.

- ModelarDB Core includes a few models, but users can load more dynamically (no need to recompile/restart)
- Models and segments must implement this interface to be used by ModelarDB

Query Examples



```
1  SELECT SUM(Value) FROM DataPoint WHERE Tid = 3
2  SELECT SUM_S(*) FROM Segment WHERE Tid = 3
3
4  SELECT AVG_SS( START(*, '2012-01-03 12:30') )
5  FROM Segment WHERE EndTime > '2012-01-03 12:30'
6
7  SELECT * FROM DataPoint WHERE Tid = 3
8  AND TS < '2012-04-22 12:25'
```

Listing 2: Query examples supported in ModelarDB

- A set of example queries supported by ModelarDB's two views
- Operations on the Segment View are implemented as **UDAFs and UDFs**

Evaluation - Storage



- 6 + 1 Laptops, **EH** (583 GiB, 100 ms), **ER** (488 GiB, 1 s), **EP** (339 GiB, 60 s)




Figure: Storage, EH




Figure: Storage, ER




Figure: Storage, EP

- ModelarDB provides better compression using model-based storage**
 - Best compression ratio for high frequency data (EH, ER) and increases with error bound
 - Average error is 0.005% (EH), 2.5% (ER) and 0.73% (EP) for a 10% error bound
 - ModelarDB degrades gracefully with more outliers

Evaluation - Adaptability




Figure: Adaptability, EH




Figure: Adaptability, ER



Figure: Adaptability, EP

- **ModelarDB chooses an appropriate model for each part of a series**
 - **Different models used** for each data set and **linear models** are used with **0%** error bound
 - The system is **extensible** and users can implement other models to increase **adaptability**

Evaluation – Ingestion and Query Processing




Figure: Ingestion, ER




Figure: Aggregate Queries, ER




Figure: Scalability (Azure), ER


- **ModelarDB has fast ingestion, fast large aggregates and linear scalability**
 - Only InfluxDB, Cassandra, and ModelarDB can answer queries while ingesting data points
 - The paper shows ModelarDB is competitive with other systems for small scale queries

ModelarDB v1: how far did we get?



- **Summary:**
 - Storing sensor data as **simple aggregates discards valuable information**
 - **Model-based compression provides multiple benefits** over simple aggregates
 - Proposed the model-based TSMS **ModelarDB** based on:
 - ◆ A **general architecture** for a **modular model-based TSMS**
 - ◆ An algorithm for **online multi-model compression** of time series
 - ◆ A set of **methods and optimizations** for a **model-based TSMS**
 - Evaluation showed that **ModelarDB hits a sweet spot** by providing:
 - ◆ **Fast ingestion**
 - ◆ **Good compression**
 - ◆ **Fast, scalable online aggregate query processing**
- But we can do **even better...**

Next step: Exploiting correlation



- **Detecting** correlation in data is an orthogonal problem
 - We let the user hint correlation
- If the time series in a group can not (no longer) be represented by a single model, ModelarDB *splits* the group
 - To respect the error bound
 - The time series can be *joined* again later

ModelarDB+ (v2) contributions




- Compression of time series groups using multiple model types, we call this type of compression **Multi-Model Group Compression (MMGC)**
- **Group Online Lossy and lossless Extensible Multi-Model (GOLEMM)**
 - **First Multi-Model Group Compression method** for time series and model types **extended** to **compress** time series **groups**
- **Primitives** for users to effectively group time series, and a method that **automatically groups** time series **using** their **metadata** as **dimensions**
- Algorithms for executing simple and **multi-dimensional aggregate queries** on models representing values from time series groups
- **ModelarDB+** a version of the **open-source** distributed model-based time series management systems ModelarDB with our methods added:
 - Available at github.com/skejserjensen/ModelarDB under Apache 2.0



Grouping Correlated Time Series



- Additional compression is achieved by **compressing time series in groups**




- Time series first **statically** grouped using metadata:
 - Time series source and dimensions
 - Dimensions** contain hierarchically organized members
- Users can indicate correlation using our **primitives**:
 - Time series sources, members in dimensions, and the distance between two sets of dimensions



Ingesting Correlated Time Series



- Models are incrementally fitted and emitted as part of **segments with metadata**:
 - Temporary Segment**: Holds an unfinished model cached in memory for low latency
 - Finalized Segment**: Holds a finished model cached in memory and persisted to disk
- Models are fitted in sequence until all would exceed the error bound:




- ModelarDB Core includes four model types, users can **optionally add** more:
 - PMC-Mean** (Constant), **Swing** (Linear), **Facebook** (Lossless), **Uncompressed** (None)

Dynamic Grouping




- The time series in a group might **temporarily not be correlated**:
- For **example**, a **temperature** sensor can be **obscured by clouds**
- This can be **efficiently detected** as the **compression ratio is reduced**
- **Dynamically splitting and merging** groups **remedies** this **problem**:



ModelarDB+ Query Processing



- **Multidimensional** aggregates, e.g., CUBE
- Denormalized user-defined **<Dimensions>**
- Data Point View
 - Interface: Tid int, TS timestamp, Value float, <Dimensions>
- Segment View
 - Interface: Tid int, StartTime timestamp, EndTime timestamp, SI int, Mid int, Parameters blob, Gaps blob, <Dimensions>
- UDAFs for aggregation on segments are suffixed with **_S**,
 - COUNT_S
- UDAFs for aggregation over time on segments are suffixed with a time interval
 - COUNT_MINUTE, MIN_HOUR, MAX_MONTH, and SUM_YEAR



ModelarDB+ Evaluation

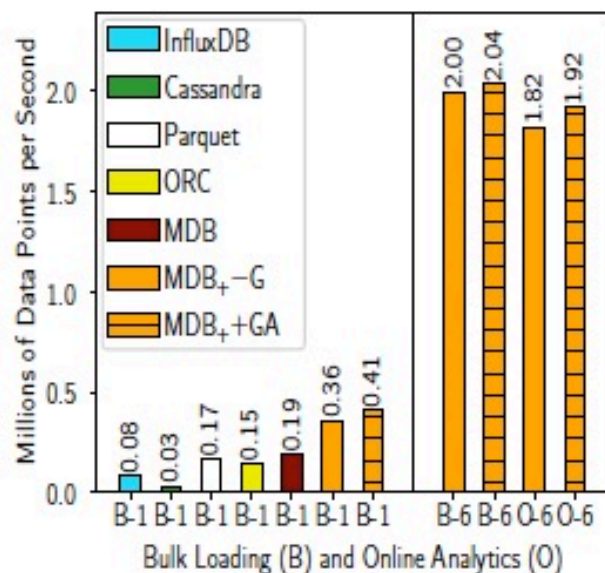


- Evaluation primarily uses **real-life data sets** from the **energy** domain:
 - **EP 45,353 time series** collected from **energy producers** with a sampling interval of **60s** and occupying **339 GiB** as uncompressed **CSV**
 - **EF 197 time series** collected from **wind turbines** with a sampling interval of **200ms** and occupying **372 GiB** as uncompressed **CSV**
- Most experiments are performed on a small cluster of commodity PCs
- Scalability experiments are performed using Microsoft Azure
- ModelarDB+ configurations:
 - with **no** grouping (MDB+ -G)
 - with **auto** (MDB+ +GA)
 - with the **best** (handtuned) primitives per data set (MDB+ +GB)

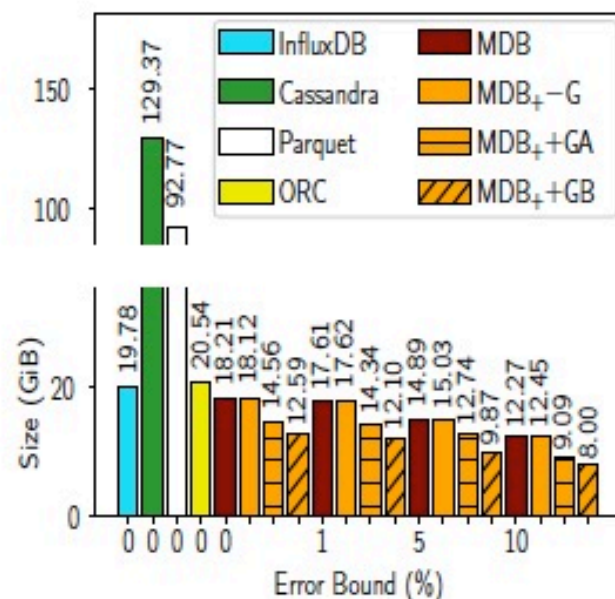
Ingestion and Storage Results ModelarDB+



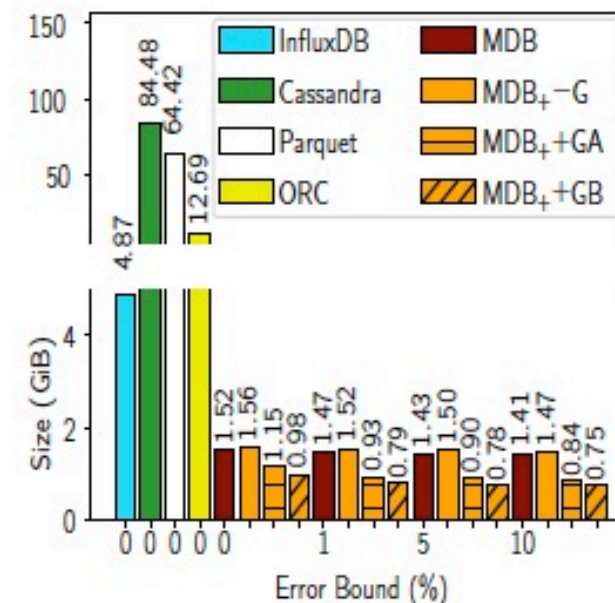
- ModelarDB+ provides both a much **higher ingestion rate** and requires much **less storage** than InfluxDB, Cassandra, Parquet, and ORC
- Using **multiple model types** allows GOLEMM to **automatically adapt**
- **Grouping improves the ingestion rate** due to the higher compression
- **Creating groups automatically** from metadata **improves the compression**



Ingestion Rate, EP



Storage Used, EP

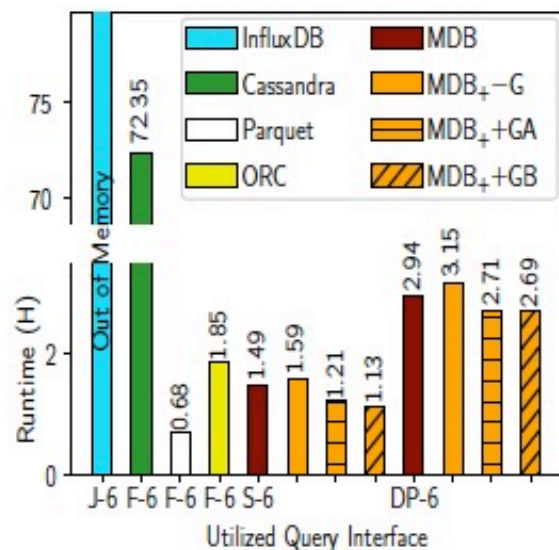


Storage Used, EH

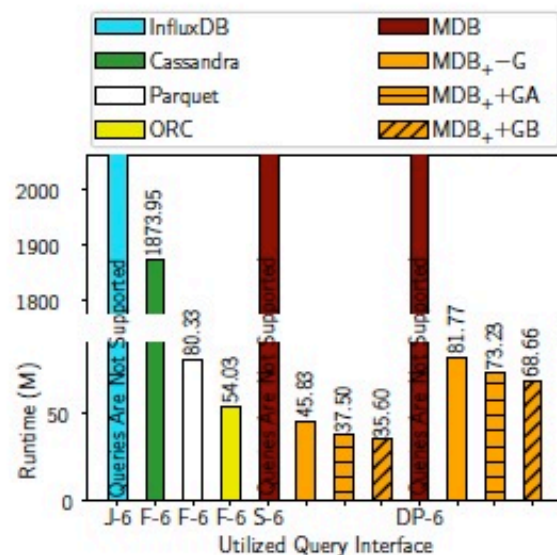
Aggregate Query Results ModelarDB+




- ModelarDB+ is **faster than** the other formats **when timestamps and values are used**, and experiments on **Azure** show that it **scales linearly**
- **Grouping time series decreases query time** for queries that **use most or all** of the time **series in each group** (Large Scale / Month and Category)
- Grouping **can also increase** query time if the query **only use a few time series from each group** (Small Scale) or if the groups are on one worker



Large Scale, EP



Month and Category, EP



Small Scale, EH

ModelarDB+ Conclusion: where are we now?



- Summary
 - **Wind turbines** produce **huge time series** that should be stored/queried at **high frequency**
 - **Current practice**: storing sensor data as **simple aggregates**, discards valuable information
 - **Grouping time series** and **storing them as models** provides **many benefits** over storing them as simple aggregates or raw data points
 - We proposed methods for creating (Primitives), compressing (the MMGC method GOLEMM), and querying time series groups
 - The evaluation of ModelarDB+ showed that **grouping** can **provide even faster ingestion** speed, **reduced storage** required, and **faster aggregate queries**
- Future Work
 - **Indexing** techniques exploiting that data is stored as models
 - **Query and cluster-aware** grouping and partitioning methods
 - Support for **high-level analytical queries and machine learning directly on models**

MORE project



- **Management Of Real-time Energy data (MORE)**
 - Call topic ICT-51-2020: Big Data technologies and extreme-scale analytics
 - October 2020 - September 2023
 - Athena RC (coordinator), AAU, InAccess, IBM Research Dublin, Perception Dynamics, LABORELEC (ENGIE subsidiary), ModelarData
- ModelarDB concept both for **edge computing** and **cloud**
 - **Optimizing edge storage** and **transfer to cloud with models**
- Advanced **time series analytics** and **machine learning directly on (streaming) models**
- Main use cases
 - **Massive solar park streams** (Inaccess)
 - **Massive wind park streams** (ENGIE/Laborelec)



ModelarData Spinout



- ModelarData spinout established for commercial exploitation and uptake
- Cloudera-like business model
 - Free open source base version
 - Paid premium feature
 - ◆ Managed subscription, documentation, training, support, consulting + custom development
- Focus: **extreme-scale analytics** mainly for **renewable energy data**

ModelarData

References



- Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, Time Series Management Systems: A Survey. In *IEEE Transactions on Knowledge and Data Engineering*, Volume 29, Number 11, Pages 2581–2600, November, 2017.
- Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, ModelarDB: Modular Model-Based Time Series Management with Spark and Cassandra. In *Proceedings of the VLDB Endowment*, Volume 11, Number 11, Pages 1688–1701, July, 2018.
- Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, Demonstration of ModelarDB: Model-Based Management of Dimensional Time Series. In *Proceedings of the 2019 International Conference on Management of Data*, Pages 1933–1936, 2019.
- Søren Kejser Jensen, Torben Bach Pedersen, Christian Thomsen, Scalable Model-Based Management of Correlated Dimensional Time Series in ModelarDB+". In the *37th IEEE International Conference on Data Engineering (ICDE)*, 2021
- MORE Project: <https://more2020.eu>
- ModelarData <https://modelardata.com>